

Assignment No. 1
Linux Command
Create File in Linux

vi firstfile

Show File Contain in Linux

cat firstfile.txt

Hadoop Command

1. Create a Directory in HDFS

To create a directory in HDFS:

hdfs dfs -mkdir firstdir

2. List Files in a Directory in HDFS

hdfs dfs -ls

3. Upload Files from Local File System to HDFS

Use the -put command to upload files from your local file system to HDFS.

hdfs dfs -put firstfile.txt /user/cloudera/firstdir

4. List Files in a Directory

Use the -ls command to list the files in an HDFS directory.

hdfs dfs -ls /user/cloudera/firstdir/

5. Display the Contents of a File

Use the -cat command to display the contents of a file in HDFS.

hdfs dfs -cat /user/cloudera/firstfile.txt

6. Delete a File in HDFS

Use the -rm command to remove a file from HDFS.

hdfs dfs -rm /user/cloudera/firstdir/firstfile1.txt

Assignment No. 2

Implement a Java program to interact with HDFS (reading and writing files).

```
import java.io.File;
```

```
public class filehand {
```

```
public static void main(String []args) throws IOException
```

```
{
```

```
File obj1=new File("/home/cloudera/Desktop/Firstfile.txt");
```

```
if (obj1.createNewFile())
```

```
System.out.print("File is Created");
```

```
else
```

```

System.out.print("File is Already exist");
FileWrite w1=new FileWriter("/home/cloudera/Desktop/Firstfile.txt");
w1.write("Welcome my first file is write");
w1.close();
Scanner r1=new Scanner(obj1);
while(r1.hasNextLine())
{
String data=r1.nextLine();
System.out.println(data);
} } }

```

Assignment No. 3

Use Hadoop's built-in commands to manage files and directories.

1. Create Directories in HDFS.

```
hdfs dfs -mkdir firstdir
```

2. Upload Files from Local File System to HDFS

Use the -put command to upload files from your local file system to HDFS.

```
hdfs dfs -put firstfile.txt /user/cloudera/firstdir
```

3. List Files in a Directory

Use the -ls command to list the files in an HDFS directory.

```
hdfs dfs -ls /user/cloudera/firstdir/
```

5. Display the Contents of a File

Use the -cat command to display the contents of a file in HDFS.

```
hdfs dfs -cat /user/cloudera/firstfile.txt
```

6. Copy Files from HDFS to Local File System

Use the -get command to copy files from HDFS to your local file system.

```
hdfs dfs -get /user/cloudera/firstdir/firstfile.txt/home/cloudera/lindir
```

7. Delete a File in HDFS

Use the -rm command to remove a file from HDFS.

```
hdfs dfs -rm /user/cloudera/firstdir/firstfile1.txt
```

7. Delete a Directory in HDFS

Use the -rm -r command to delete a directory and its contents from HDFS.

```
hdfs dfs -rm -r /user/cloudera/firstdir (For Non Empty Directory
```

```
hdfs dfs -rm /user/cloudera/firstdir (For Empty Directory )
```

Assignment No. 4

Implement Map Side Join and Reduce Side Join.

(Write hadoop code to implement Map Reduce application count number of

word in file)

```
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WordCount {
    public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }

    public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();
        public void reduce(Text key, Iterable<IntWritable> values,
        Context context

        ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }
}
```

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

File Link

https://hadoop.apache.org/docs/current/hadoop-mapreduce-client/hadoop-mapreduce-client-core/MapReduceTutorial.html#Source_Code

Step 1: Export Java Eclipse Project Jar File to Cloudera

Step 2. Make firstfile.txt file vi editor ->Write data

Step 3: Perform Below commands on terminal

Command Map Reduce Code

1) Transfer all local file to hadoop

Hdfs dfs -put firstfile.txt /user/cloudera

Hdfs dfs -put WordCount.jar /user/cloudera

2) Run Java Jar File for Map Reduce Operation

hadoop jar WordCount.jar WordCount firstfile.txt outputfile

3) List outputfile

hdfs dfs -ls /user/cloudera/outputfile

4) Show outputfile

hdfs dfs -cat /user/cloudera/outputfile/part-r-00000

Assignment No. 6

Pipeline multiple Map Reduce jobs

Example: Pipelining Two Jobs

Job 1: Word Count (Word frequency count)

This first job counts the occurrences of each word in the input text files.

```
import java.io.IOException; import
java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable; import
org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job; import
org.apache.hadoop.mapreduce.Mapper; import
org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat; import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
public class WordCount {
    public static class TokenizerMapper extends Mapper<Object, Text, Text,
        IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context) throws
            IOException, InterruptedException {
            String[] words = value.toString().split("\\s+"); for
            (String wordStr : words) {
                word.set(wordStr);
                context.write(word, one);
            }
        }
    }
    public static class IntSumReducer extends Reducer<Text, IntWritable, Text,
        IntWritable> {
        private IntWritable result = new IntWritable();
        public void reduce(Text key, Iterable<IntWritable> values, Context
            context) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum); context.write(key,
```

```

result);
}
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Job 2:Filter Words with Frequency GreaterThan 2

The second job processes the output of the first job to filter and only output words that have a frequency greater than 2

```

public class FilterWords {
    public static class FilterMapper extends Mapper<Object, Text, Text,
    IntWritable> {
        private IntWritable count = new IntWritable();
        public void map(Object key, Text value, Context context) throws
        IOException, InterruptedException {

```

```

        String[] fields = value.toString().split("\\t");
        String word = fields[0];
        int wordCount = Integer.parseInt(fields[1]);
        // Output only words with count greater than 2 if
        (wordCount > 2) {
            count.set(wordCount); context.write(new
            Text(word), count);
        }
    }
}

```

```

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "filter words");

```

```

job.setJarByClass(FilterWords.class);
job.setMapperClass(FilterMapper.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0])); // Input path
from the first job's output
FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output
path
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

☐ Output format: The output of each job should be compatible with the input format of the next

job. For instance, if the output of Job 1 is a simple key-value pair (word and count), Job 2 should be able to process that format directly.

Step 1: Export Java Eclipse Project Jar File to Cloudera Step 2. Make firstfile.txt file vi editor -> Write data Step 3: Perform Below commands on terminal

Command MapReduce Code

1) Transfer all local file to hadoop

Hdfs dfs -put firstfile.txt /user/cloudera Hdfs

dfs -put PipLine1.jar /user/cloudera

2) Run First job of Java Jar File for Map Reduce Operation

hadoop jar PipLine1.jar wordcount firstfile.txt outpip1

3) Run Second job of Java Jar File for Map Reduce Operation

hadoop jar PipLine1.jar FilterWords outpip1 outpip2

4) List outputfile

hdfs dfs -ls /user/cloudera/outpip2

5) Show outputfile

hdfs dfs -cat /user/cloudera/outpip2/part-r-00000

Assignment No. 7

Create and use UDFs in Pig Latin scripts (Write hadoop code to convert username in uppercase)

Step 1: Open Java Eclipse -> Make New Project -> (PigAss7) Add External Library -> pig , hadoop

,hadoop 0.20_map_reduce -> finish

Step 2: Add class in project -> PigUDF -> Write below code (Create the UDF)

```

import org.apache.pig.EvalFunc;
import org.apache.pig.data.Tuple;

```

```

public class PigUDF extends EvalFunc<String> {
public String exec(Tuple tuple)throws IOException {
if(tuple ==null) {
return null;
}
String user=(String)tuple.get(0);
String city=(String)tuple.get(1);
int score=(Integer)tuple.get(2);
return user+","+city.toUpperCase()+" "+score;
}
}

```

Step 3: Export java project in to jar file ->PigAss7.jar

Step 4: Open terminal create cust_us.txt file using following command

vi cust_us

ex. User1,dhule,3

Step 5: Open terminal to create pig file using following command.

vi toupper.pig

Write following code in file

REGISTER /home/cloudera/PigAss7.jar

DEFINE toupper PigUDF();

usa = LOAD '/home/cloudera/cust_us.txt' Using PigStorage(',')

AS (user:chararray,city:chararray,score:int);

usa_upper = FOREACH usa GENERATE toupper(user,city,score);

DUMP usa_upper;

Step 6: Run Pig script

pig -x local toupper.pig

Assignment No. 7

Create and use UDFs in Pig Latin scripts (Write hadoop code to convert username in uppercase)

Step 1: Open Java Eclipse -> Make New Project->(PigAss7) Add External Library->pig ,
hadoop

,hadoop 0.20_map_reduce->finish

Step 2: Add class in project -> PigUDF -> Write below code (Create the UDF)

```
import org.apache.pig.EvalFunc;
```

```
import org.apache.pig.data.Tuple;
```

```
public class PigUDF extends EvalFunc<String> {
```

```
public String exec(Tuple tuple)throws IOException {
```



```

if(tuple ==null) {
return null;
}
String user=(String)tuple.get(0);
String city=(String)tuple.get(1);
int score=(Integer)tuple.get(2);
return user+","+city.toUpperCase()+","+score;
}
}

```

Step 3: Export java project in to jar file ->PigAss7.jar

Step 4: Open terminal create cust_us.txt file using following command

vi cust_us

Step 5: Open terminal to create pig file using following command.

vi toupper.pig

Write following code in file

REGISTER /home/cloudera/PigAss7.jar

DEFINE toupper PigUDF();

usa = LOAD '/home/cloudera/cust_us.txt' Using PigStorage(',')

AS (user:chararray,city:chararray,score:int);

usa_upper = FOREACH usa GENERATE toupper(user,city,score);

DUMP usa_upper;

Step 6: Run Pig script

pig -x local toupper.pig

Assignment No 9

Implement and execute HiveQL queries to perform data retrieval and manipulation.

1. Setting up Hive

Hive

2. Creating a Hive Table

CREATE TABLE employees (

emp_id INT,

emp_name STRING,

emp_salary DOUBLE,

emp_department STRING

) ROW FORMAT DELIMITED

FIELDS TERMINATED BY ','

STORED AS TEXTFILE;

3. Data Manipulation

a. Inserting Data into the Table:

```
INSERT INTO TABLE employees VALUES (1001, 'John Doe', 60000.00, 'Engineering');
```

b. Updating Data:

```
INSERT OVERWRITE TABLE employees
SELECT emp_id, emp_name, CASE
WHEN emp_salary < 50000 THEN emp_salary * 1.1
ELSE emp_salary
END AS emp_salary,
emp_department
FROM employees;
```

c. Deleting Data:

```
INSERT OVERWRITE TABLE employees
SELECT * FROM employees WHERE emp_id != 1001;
```

4. Basic Data Retrieval Queries

a. Select all records:

```
SELECT * FROM employees;
```

b. Select specific columns:

```
SELECT emp_name, emp_salary FROM employees;
```

c. Select with a WHERE clause:

```
SELECT * FROM employees WHERE emp_salary > 50000;
```

Assignment No 10

Perform operations like joins, group by, and aggregations in Hive.

1. Creating a Hive Table

```
CREATE TABLE employees (
emp_id INT,
emp_name STRING,
emp_salary DOUBLE,
emp_department STRING
) ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;
CREATE TABLE departments (
dept_id INT,
dept_name STRING,
) ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS TEXTFILE;
```

Insert 5 record on each table by using insert command

2. Joins

```
SELECT e.emp_name, d.dept_name
```

```
FROM employees e  
JOIN departments d  
ON e.emp_department = d.dept_id;
```

3. Aggregate Functions

a. Count the number of employees:

```
SELECT COUNT(*) FROM employees;
```

b. Find the average salary:

```
SELECT AVG(emp_salary) FROM employees;
```

c. Find the highest salary:

```
SELECT MAX(emp_salary) FROM employees;
```

4. Group By

Group by department and get average salary per department:

```
SELECT emp_department, AVG(emp_salary) FROM employees GROUP BY emp_department;
```